

```
!pip install -qU gradio langchain-cohere faiss-cpu python-dotenv langchain-community
```

```
import os
from langchain_core.example_selectors import SemanticSimilarityExampleSelector
from langchain_core.prompts import FewShotPromptTemplate, PromptTemplate
from langchain_community.vectorstores import FAISS
```

✓ Config and Params

```
MODEL_TEMPERATURE = 0.1
TOP_N_MATCH = 2 # Tune it as per number of examples in input data

# Load from .env file
os.environ["COHERE_API_KEY"] = "xxxx"

DEFAULT_ANSWER = '"Sorry, I can not provide answer to given question."'

SYSTEM_PROMPT = f""""You are an enterprise grader customer support agent for Thoughtful AI.

# TASK
- You are provided some examples and a user question.
- You need to answer user question only as per given example match.

# RULES
- If user question matches any example, provide the answer as it is mentioned in the example word by word.
- If user question doesn't match any example, provide MANDATORILY {DEFAULT_ANSWER}.
- Never answer any user question which does not match with provided examples.
- If there is no match, only say {DEFAULT_ANSWER}.
"""

Start coding or generate with AI.
```

✓ Loading Models

```
from langchain_cohere import CohereEmbeddings, ChatCohere

if "COHERE_API_KEY" not in os.environ:
    raise ValueError("COHERE_API_KEY environment variable not set. Please set your Google API key.")

embeddings = CohereEmbeddings(
    model="embed-english-v3.0",
    # input_type="search_query" # Optimized for question answering
)

llm = ChatCohere(
    model="command-r",
    temperature=MODEL_TEMPERATURE,
    max_tokens=1024
)
```

✓ Raw Data

```
thoughtful_ai_qa = {
    "questions": [
        {
            "question": "What does the eligibility verification agent (EVA) do?",
            "answer": "EVA automates the process of verifying a patient's eligibility and benefits information in real-time, eliminating manual intervention, and improving accuracy."
        },
        {
            "question": "What does the claims processing agent (CAM) do?",
            "answer": "CAM streamlines the submission and management of claims, improving accuracy, reducing manual intervention, and accelerating payment processing."
        },
        {
            "question": "How does the payment posting agent (PHIL) work?",
            "answer": "PHIL automates the posting of payments to patient accounts, ensuring fast, accurate reconciliation of payments and reducing manual intervention."
        },
        {
            "question": "Tell me about Thoughtful AI's Agents.",
            "answer": "Thoughtful AI provides a suite of AI-powered automation agents designed to streamline healthcare processes. These agents include the Eligibility Verification Agent (EVA), Claims Processing Agent (CAM), and Payment Posting Agent (PHIL), each designed to improve efficiency and accuracy while reducing manual intervention."
        }
    ]
}
```

```

        "question": "What are the benefits of using Thoughtful AI's agents?",
        "answer": "Using Thoughtful AI's Agents can significantly reduce administrative costs, improve operational efficiency, and r
    }
]
}

examples = [
    {"input": q["question"], "output": q["answer"]}
    for q in thoughtful_ai_qa["questions"]
]

print(f"Number of examples: {len(examples)}")
examples[0]

```

➡ Number of examples: 5

```

{'input': 'What does the eligibility verification agent (EVA) do?',
 'output': 'EVA automates the process of verifying a patient's eligibility and benefits information in real-time, eliminating
manual data entry errors and reducing claim rejections.'}

```

✓ Vectorization

✓ Key Implementation Details:

1. Vector Similarity Search:

- Uses `SemanticSimilarityExampleSelector` with FAISS vector store
- Embeds questions using OpenAI's text-embedding-3-small model
- Retrieves top 3 similar questions

2. Few-Shot Prompting:

- Combines retrieved examples with the user's question
- Uses structured prompt template to guide the LLM

Start coding or [generate](#) with AI.

```

example_selector = SemanticSimilarityExampleSelector.from_examples(
    examples,
    embeddings,
    FAISS,
    k=TOP_N_MATCH # Retrieve top N similar examples
)

```

```

# %% Create few-shot prompt template
example_prompt = PromptTemplate(
    input_variables=["input", "output"],
    template="Question: {input}\nAnswer: {output}"
)

```

✓ Retrieve Top-N

```

few_shot_prompt = FewShotPromptTemplate(
    example_selector=example_selector,
    example_prompt=example_prompt,
    prefix="You are a customer support agent for Thoughtful AI. Answer questions using these examples:",
    suffix="Question: {input}\nAnswer:",
    input_variables=["input"]
)

```

✓ Chatbot Agent

✓ 1. Semantic Retrieval from Predefined Dataset

- The agent uses semantic similarity search to retrieve the most relevant answers from a hardcoded dataset about Thoughtful AI.
- By leveraging vector embeddings and FAISS, it ensures that user queries are matched to the closest available knowledge, even if the wording differs.

2. Few-Shot Prompt Construction

- For each user question, the agent dynamically constructs a prompt using the top-matching examples from the dataset.
- This few-shot approach improves answer accuracy and relevance, as the LLM is guided by real, contextually similar Q&A pairs.

3. LLM Fallback for Out-of-Domain Queries

- If the agent cannot find any relevant examples in the dataset, it automatically falls back to a generic LLM response.
- This ensures that the chatbot remains conversational and helpful, even for queries outside the predefined scope.

4. Robust Error Handling

- The function is wrapped in a try-except block to catch and report any unexpected errors gracefully.
- Handles missing or malformed data in the examples, skipping problematic entries without interrupting the user experience.
- If no relevant examples are found, the agent still provides a meaningful response rather than failing silently.

5. Maintainability and Extensibility

- The function is modular and easy to update: new Q&A pairs can be added to the dataset without code changes.
- The agent logic is separated from the UI layer, making it adaptable for CLI, web, or API-based interfaces.

6. Session History Support

- The function signature includes a `history` parameter, allowing for future enhancements such as context-aware multi-turn conversations.

```
# Main Chat Agent
def chatbot_ai_agent(message: str, history: list = None) -> str:
    try:
        # Retrieve similar examples (modified to handle empty results)
        similar_examples = example_selector.select_examples({"input": message}) or []

        if not similar_examples:
            return llm.invoke(f"Answer about Thoughtful AI: {message}").content

        # Format examples with proper error handling
        formatted_examples = []
        for ex in similar_examples:
            try:
                formatted_examples.append(example_prompt.format(**ex))
            except KeyError:
                continue

        # Build final prompt with fallback
        final_prompt = f"""\{SYSTEM_PROMPT}\.

        # Examples:
        {''.join(formatted_examples)}

        # User question:
        {message}
        Answer:"""

        return llm.invoke(final_prompt).content

    except Exception as e:
        return f"Sorry, Something went wrong..!!"
```

✓ Testing

Case-1: Exact Match

chatbot_ai_agent(message="What does the eligibility verification agent (EVA) do?")

↳ 'EVA automates the process of verifying a patient's eligibility and benefits information in real-time, eliminating manual data entry.'

Case-2: Semantic Match

chatbot_ai_agent(message="What does EVA do?")

↳ 'EVA automates the process of verifying a patient's eligibility and benefits information in real-time, eliminating manual data entry.'

Case-3: No Match

chatbot_ai_agent(message="what is capital of USA?")


↳ 'Sorry, I can not provide answer to given question.'

✓ Chatbot

```
import gradio as gr
```

```
# Gradio interface with validated parameters
```


```
demo = gr.ChatInterface(
    fn=chatbot_ai_agent,
    title="Thoughtful AI Support Agent",
    description="Ask about Thoughtful AI's products and solutions",
    examples=[q["question"] for q in thoughtful_ai_qa["questions"]],
    cache_examples=True,
    autofocus=True,
    multimodal=False # Explicitly disable multimodal input
)
```

 /usr/local/lib/python3.11/dist-packages/gradio/chat_interface.py:338: UserWarning: The 'tuples' format for chatbot messages is deprecated.
self.chatbot = Chatbot()

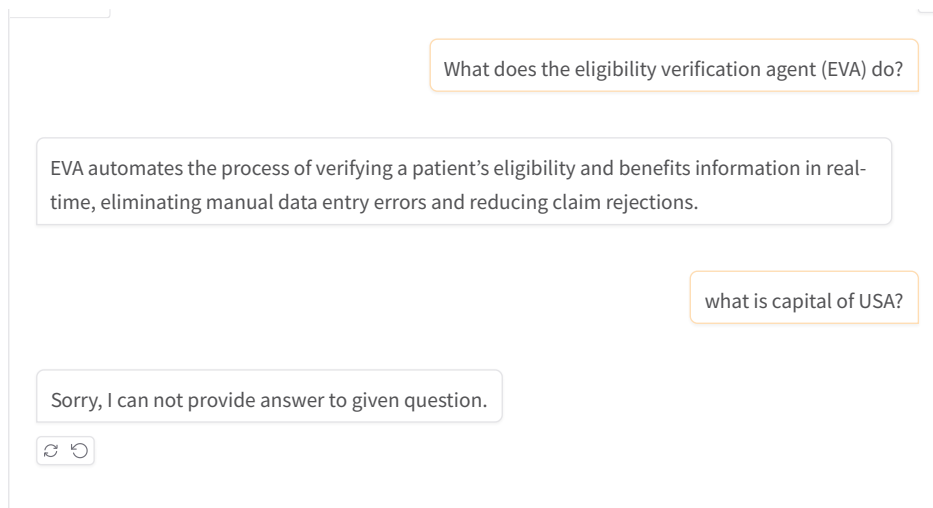
Start coding or [generate](#) with AI.

```
# Final launch
```

```
demo.launch(
    share=True,
    inline=True,
    show_error=True,
    debug=False # Disable debug mode for production
)
```

 Caching examples at: '/content/.gradio/cached_examples/17'
Colab notebook detected. To show errors in colab notebook, set debug=True in launch()
* Running on public URL: <https://66636bdafe9831883e.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working



Start coding or [generate](#) with AI.

